

Cost ratio per GPU, deep columns only ($\text{nz} \geq 50$)

linear scale: by $\text{nz}=50$ the MI250X frame penalty is gone
and all four agree to within a few percent

