

Cost ratio per GPU, deep columns only ($nz \geq 50$)

linear scale: by $nz=50$ the MI250X frame penalty is gone
and all four agree to within a few percent

