

Cost ratio per GPU, deep columns only (nz ≥ 50)

473×297 production grid, linear scale: by nz=50 the MI250X frame penalty is gone and all four agree to within a few percent

cost ratio = $t(\text{NZ_STACK_MAX}=128) \div t(\text{NZ_STACK_MAX}=\text{nz}+1)$

V100 GH200 MI250X Intel Max

cost ratio, production ÷ fitted

